



The Relational Model

Abdu Alawini

University of Illinois at Urbana-Champaign

CS411: Database Systems

Learning Objectives

After this lecture, you should be able to:

- Define a **data model**
- Define the **relational data model**
- Articulate the **basic terminologies** of the **relational data model** (from a practical perspective)
- Define **Primary and Foreign keys**

What is a Data Model?

A data model is a notation for **describing data or information**. The description generally consists of three parts:

1. *Structure of the data:*

- data structures used to implement data in the computer (physical data model)

2. *Operations on the data:*

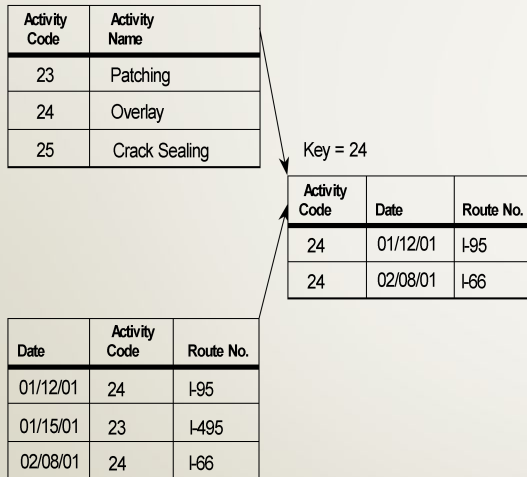
- limited set of queries (operations that retrieve information) and modifications (operations that change the database).

3. *Constraints on the data:*

- ways to describe limitations on what the data can be. These constraints often come from the real-world application requirements

Data Model Examples

Relational Model



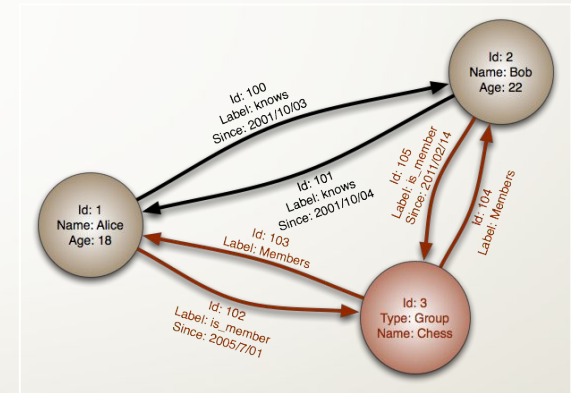
Src: Wikipedia

Document (e.g., JSON)

```
{
  "first name": "John",
  "last name": "Smith",
  "age": 25,
  "address": {
    "street address": "21 2nd Street",
    "city": "New York",
    "state": "NY",
    "postal code": "10021"
  },
  "phone numbers": [
    {
      "type": "home",
      "number": "212 555-1234"
    },
    {
      "type": "fax",
      "number": "646 555-4567"
    }
  ]
}
```

Src: Wikipedia

Graph Model



Src: Wikipedia

Outline

Data Models

- Relational Database Model
 - Basic Concepts and Terminology
 - Keys and Foreign Keys
 - Schema Specifications


The Relational Data Model

It all began with a breakthrough paper by E.F. Codd in 1970:
“A relational model of data for large shared data banks”.

Communications of the ACM 13 (6): 377

Codd's insights:

- Separate physical implementation from logical.
- Describe the data and operations **mathematically**.



Database from a user Perspective

- We'll assume that a DB has been already been implemented and loaded with data
- Our roles is to query/modify the data using SQL
- But before that, we need to learn the basics of relational model (from a practical point of view)

Introduction to Relational Databases from a Practical Point of View

Account

Number	Owner	Balance	Type
101	J. Smith	1000.00	checking
102	W. Wei	2000.00	checking
103	J. Smith	5000.00	savings
104	M. Jones	1000.00	checking
105	H. Martin	10,000.00	checking

Imagine that this table (or relation) has been defined to help keep track of bank accounts.

Table Structure

The *name* of the table

Account

The name of the *columns* (*attributes*)

Number	Owner	Balance	Type
101	J. Smith	1000.00	checking
102	W. Wei	2000.00	checking
103	J. Smith	5000.00	savings
104	M. Jones	1000.00	checking
105	H. Martin	10,000.00	checking

Table Schema

The *schema* for the table

Account			
Number	Owner	Balance	Type
101	J. Smith	1000.00	checking
102	W. Wei	2000.00	checking
103	J. Smith	5000.00	savings
104	M. Jones	1000.00	checking
105	H. Martin	10,000.00	checking

The **schema** sets the structure of the table. You can think of the schema as the *definition* of the table. (Note, the schema specifies more information than what is shown.)

Table Rows

Account

Number	Owner	Balance	Type
101	J. Smith	1000.00	checking
102	W. Wei	2000.00	checking
⋮	J. Smith	5000.00	savings
104	M. Jones	1000.00	checking
105	H. Martin	10,000.00	checking

Each entry in the table is called a *row* (*tuple*).

Sometimes an entry in the table is called a record.

Table Instance

An *instance* of the table...

the current contents or data in the table.

Account

Number	Owner	Balance	Type
101	J. Smith	1000.00	checking
102	W. Wei	2000.00	checking
103	J. Smith	5000.00	savings
104	M. Jones	1000.00	checking
105	H. Martin	10,000.00	checking

Another Table Instance

Another *instance* of the table
(two rows added, one (103) deleted)

Account

Number	Owner	Balance	Type
101	J. Smith	1,000.00	checking
102	W. Wei	2,000.00	checking
104	M. Jones	1,000.00	checking
105	H. Martin	10,000.00	checking
107	W. Yu	7,500.00	savings
109	R. Jones	432.55	checking

new

Intension vs. Extension

The *intension* of the table

Account



The diagram illustrates the concepts of intension and extension for a table. A red box highlights the table's header, labeled 'The intension of the table'. A green box highlights the table's data rows, with a green arrow pointing to it from the text 'The extension of the table. Also called the extent.' below.

Number	Owner	Balance	Type
101	J. Smith	1000.00	checking
102	W. Wei	2000.00	checking
103	J. Smith	5000.00	savings
104	M. Jones	1000.00	checking
105	H. Martin	10,000.00	checking


The *extension* of the table. Also called the *extent*.

“Size” of a Table

Degree or arity of a table is the number of columns

Degree of this relation (or table) is 4
because there are 4 attributes

Account



Number	Owner	Balance	Type
101	J. Smith	1000.00	checking
102	W. Wei	2000.00	checking
103	J. Smith	5000.00	savings
104	M. Jones	1000.00	checking
105	H. Martin	10,000.00	checking

Cardinality
of this instance
is 5 (because
there are 5
rows)

Cardinality of a table = the number of rows in the
current instance

Outline

✓ Data Models

- Relational Database Model

✓ Basic Concepts and Terminology

- Keys and Foreign Keys
- Schema Specifications

Database (One or More Tables)

Account	Number	Owner	Balance	Type	
	101	J. Smith	1000.00	checking	
	102	W. Wei	2000.00	checking	
	103	J. Smith	5000.00	savings	
	104	M. Jones	1000.00	checking	
	105	H. Martin	10,000.00	checking	

Deposit	AcctNo	Transaction-id	Date	Amount	
	102	1	10/22/00	500.00	
	102	2	10/29/00	200.00	
	104	3	10/29/00	1000.00	
	105	4	11/02/00	10,000.00	

Check	AcctNo	Check-number	Date	Amount	
	101	924	10/23/00	125.00	
	101	925	10/24/00	23.98	

Table Keys

Account	Number	Owner	Balance	Type	
	101	J. Smith	1000.00	checking	
	102	W. Wei	2000.00	checking	
	103	J. Smith	5000.00	savings	
	104	M. Jones	1000.00	checking	
	105	H. Martin	10,000.00	checking	

Deposit	AcctNo	Transaction-id	Date	Amount
	102	1	10/22/00	500.00
	102	2	10/29/00	200.00
	104	3	10/29/00	1000.00
	105	4	11/02/00	10,000.00

Check	AcctNo	Check-number	Date	Amount
	101	924	10/23/00	125.00
	101	925	10/24/00	23.98

Each table has a key.... where the values must be unique.

Table Keys (cont.)

Account	Number	Owner	Balance	Type	
	101	J. Smith	1000.00	checking	
	102	W. Wei	2000.00	checking	
	103	J. Smith	5000.00	savings	
	104	M. Jones	1000.00	checking	
	105	H. Martin	10,000.00	checking	

Deposit	AcctNo	Transaction-id	Date	Amount
	102	1	10/22/00	500.00
	102	2	10/29/00	200.00
	104	3	10/29/00	1000.00
	105	4	11/02/00.	10,000.00

Check	AcctNo	Check-number	Date	Amount
	101	924	10/23/00	125.00
	101	925	10/24/00	23.98

Key may consist of one column or two (or more) columns.

Connections between Tables

Account	Number	Owner	Balance	Type	
	101	J. Smith	1000.00	checking	
	102	W. Wei	2000.00	checking	
	103	J. Smith	5000.00	savings	
	104	M. Jones	1000.00	checking	
	105	H. Martin	10,000.00	checking	

Deposit	AcctNo	Transaction-id	Date	Amount
	102	1	10/22/00	500.00
	102	2	10/29/00	200.00
	104	3	10/29/00	1000.00
	105	4	11/02/00	10,000.00
	106	5	12/05/00	555.00

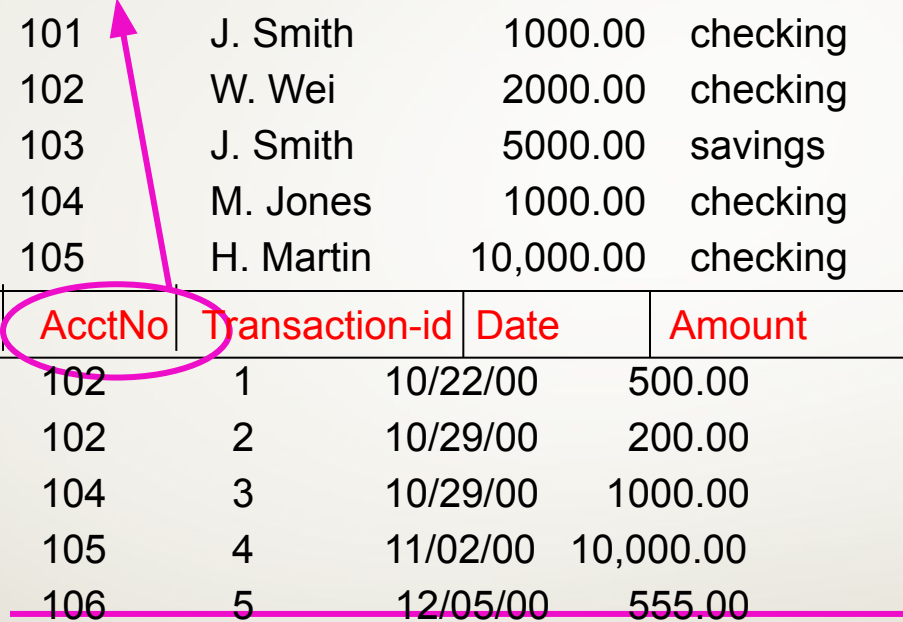


Is this legal?

If not, how do we prevent it from happening?

Foreign Key

Account	Number	Owner	Balance	Type	
	101	J. Smith	1000.00	checking	
	102	W. Wei	2000.00	checking	
	103	J. Smith	5000.00	savings	
	104	M. Jones	1000.00	checking	
	105	H. Martin	10,000.00	checking	
Deposit	AcctNo	Transaction-id	Date	Amount	
	102	1	10/22/00	500.00	
	102	2	10/29/00	200.00	
	104	3	10/29/00	1000.00	
	105	4	11/02/00	10,000.00	
	106	5	12/05/00	555.00	



We say that **Deposit.AcctNo** is a *foreign key* that *references* **Account.Number**. If the DBMS enforces this constraint, we have *referential integrity*.

Foreign keys might or might not be part of the key for the referring table

Account	Number	Owner	Balance	Type	
	101	J. Smith	1000.00	checking	
	102	W. Wei	2000.00	checking	
	103	J. Smith	5000.00	savings	
	104	M. Jones	1000.00	checking	
	105	H. Martin	10,000.00	checking	

Deposit	AcctNo	Transaction-id	Date	Amount
Deposit.AcctNo	102	1	10/22/00	500.00
is not part	102	2	10/29/00	200.00
of key for	104	3	10/29/00	1000.00
Deposit.	105	4	11/02/00	10,000.00

Check	AcctNo	Check-number	Date	Amount
Check.AcctNo	101	924	10/23/00	125.00
is part of	101	925	10/24/00	23.98
key for Check.				

Foreign
Key

Primary
Key

Outline

✓ Data Models


- Relational Database Model

✓ Basic Concepts and Terminology

✓ Keys and Foreign Keys

- Schema Specifications

Specification of a Database Schema

- Select the tables, with a **name for each table**.
- Select **columns for each table** and give the **domain for each column**.
- Specify the **key(s)** for each table.

There can be more than one key for a table.
- Specify all appropriate **foreign keys**.

Database Domains for Columns

Account	Number	Owner	Balance	Type	
	101	J. Smith	1000.00	checking	
	102	W. Wei	2000.00	checking	
	...				

For every column of every table, **the schema specifies allowable values**. For example,

Number must be a 3-digit number

Owner must be a 30-character string

Type must be “checking” or “savings”

The set of allowable values for a column is called the **domain** of the column.

Example Database Schema

(Keys are underlined. Each table has one key.)

Student	Takes	Course									
<table><tr><td><u>sid</u></td><td>name</td></tr></table>	<u>sid</u>	name	<table><tr><td><u>sid</u></td><td>exp-grade</td><td><u>cid</u></td><td><u>sem</u></td></tr></table>	<u>sid</u>	exp-grade	<u>cid</u>	<u>sem</u>	<table><tr><td><u>cid</u></td><td>subj</td><td><u>sem</u></td></tr></table>	<u>cid</u>	subj	<u>sem</u>
<u>sid</u>	name										
<u>sid</u>	exp-grade	<u>cid</u>	<u>sem</u>								
<u>cid</u>	subj	<u>sem</u>									
	Professor	Teaches									
	<table><tr><td><u>fid</u></td><td>name</td></tr></table>	<u>fid</u>	name	<table><tr><td><u>fid</u></td><td><u>cid</u></td><td><u>sem</u></td></tr></table>	<u>fid</u>	<u>cid</u>	<u>sem</u>				
<u>fid</u>	name										
<u>fid</u>	<u>cid</u>	<u>sem</u>									

In relational DBs, we use *relation(attribute:domain)*

STUDENT(sid:int, name:string)

Takes(sid:int, exp-grade:char[2], cid:string, sem:char[3])

COURSE(cid:string, subj:string, sem:char[3])

Teaches(fid:int, cid:string, sem:char[3])

PROFESSOR(fid:int, name:string)

Popularity of the Relational Data Model

- Most popular database systems use the relational model.
 - Oracle
 - MS SQL Server
 - MySQL
 - PostgreSQL
 - IBM DB2
 - SQLite
 - Microsoft Access
- Check: <https://db-engines.com/en/ranking>
- Learning about the relational model (and SQL) is a wise investment.

Outline

✓ Data Models

- Relational Database Model

✓ Basic Concepts and Terminology

✓ Keys and Foreign Keys

✓ Schema Specifications