



Storage & Indexing

Abdu Alawini

University of Illinois at Urbana-Champaign

CS411: Database Systems

Learning Objectives

After this lecture, you should be able to:

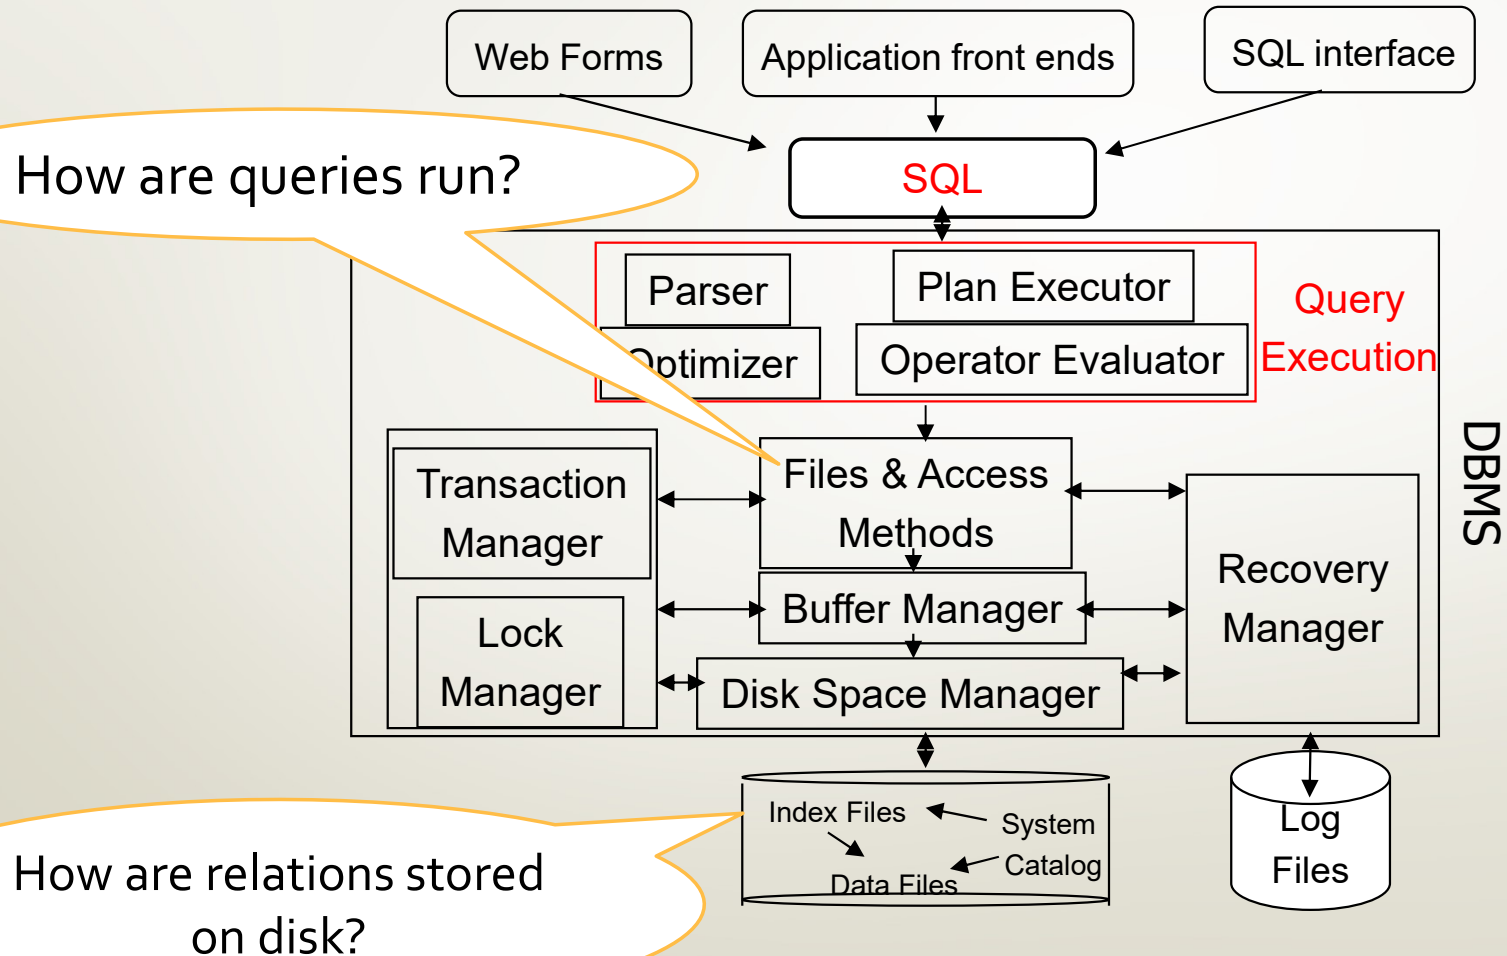
- describe how relations are stored on disk
- describe how index structures speed up data access

CS411 Goals:

Two Perspectives of DBMS

- USER PERSPECTIVE
 - **how to use a database system?**
 - conceptual data modeling, the relational and other data models, database schema design, relational algebra, SQL and No-SQL query languages.
- SYSTEMS PERSPECTIVE
 - **how to design and implement a database system?**
 - data representation, indexing, query optimization and processing, transaction processing, and concurrency control.
 - NOT COMPLETE: high-level view of implementation; CS511


DBMS Architecture




Today's lecture

- Storage
- Indexing
 - What is an index? Why do we need it?

Simplified Computer Architecture



Cost of Accessing Data on Disk



Main
Memory


Speed

10^7 ns

10^5 ns

10^2 ns

Block size vs. record size



OK. So how do we do simple operations?

Lookups. Insertions. Deletions

Outline

✓ Storage

• Indexing

- What is an index? Why do we need it?



One catch!

Indexes in databases

- An index speeds up selections on the *search key field(s)*
- Search key = any subset of the fields of a relation
 - *Search key* is **not** necessarily the same as a *key*
- Entries in an index: (k, r), where:
 - k = the search key
 - r = the record OR record id OR record ids OR pointers

Some terminology

- ***Data file***: has the data corresponding to a relation
- ***Index file***: has the index
- File consists of smaller units called **blocks** (e.g. of size 4 KB or 8 KB)
- # index blocks < # data blocks.
Index may even fit into main memory.

An Index is a Function!

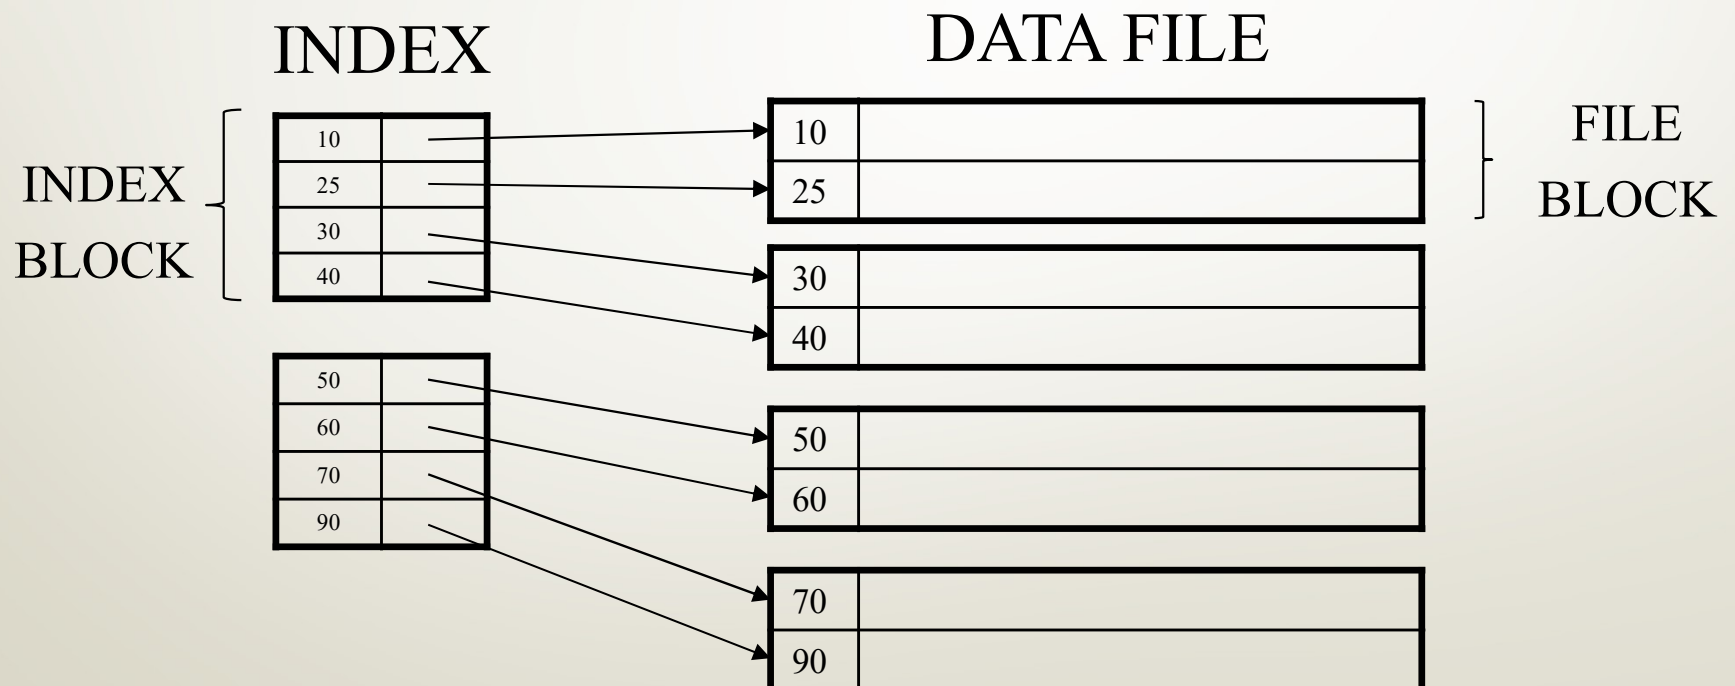
$f(\text{what: key}) = \text{where: file block}$

Characteristics of Indexes

- **Clustered/unclustered**
 - Clustered: records sorted in the search key order
 - Unclustered: records are NOT sorted in the search key order
- **Dense/sparse**
 - Dense = each record has an entry in the index
 - Sparse = only some records have
- **Primary/secondary**
 - Primary = on the primary key
 - Secondary = on any attribute


Ex: Clustered, Dense Index

- Clustered: File is sorted on the index attribute
- Dense: sequence of (key,pointer) pairs



Clustered, Sparse Index


- Sparse index: one key per data block, corresponding to the lowest search key in that block



What if there are duplicate keys?


Clustered Index with Duplicate Keys

Dense index: point to the first record with that key
(must have a pointer for each new key)



Clustered Index with Duplicate Keys


- Sparse index: pointer to lowest search key in each block



Search for 20

Clustered Index with Duplicate Keys


- *Better: pointer to lowest new search key in each block*



Search for 20 again

Unclustered Indexes

- Often for indexing other attributes than primary key
- Can it be sparse?



Summary Clustered vs. Unclustered Index

