



Indexing: Hash Tables

Abdu Alawini

University of Illinois at Urbana-Champaign

CS411: Database Systems

Learning Objectives

After this lecture, you should be able to:

- Describe how to search, insert and delete keys from
 - Secondary storage Hash Table (HT)
 - Extensible HT
 - Linear HT

Hash Tables

- Secondary storage hash tables are much like main memory ones
- Recall basics:
 - There are B buckets
 - A hash function $h(k)$ maps a key k to $\{0, 1, \dots, B-1\}$
 - Store in bucket $h(k)$ a pointer to record with key k
- Secondary storage: bucket = block
 - Store in the block of bucket $h(k)$ any record with key k
 - use overflow blocks when needed

Hash Table Example

- Assume 1 bucket (block) stores 2 records
- $h(e)=0$
- $h(b)=h(f)=1$
- $h(g)=2$
- $h(a)=h(c)=3$

0	e
1	b
2	f
3	g
	a
	c

Searching in a Hash Table

- Search for a:
- Compute $h(a)=3$
- Read bucket (block) 3
- 1 disk access

Main memory may have an array of pointers (to buckets) accessible by bucket number.

0	e
1	b f
2	g
3	a c


Insertion in Hash Table

- Place in right bucket (block), if space
- E.g. $h(d)=2$

0	e
1	b f
2	g d
3	a c

Insertion in Hash Table

- Create overflow block, if no space
- E.g. $h(k)=1$



More over-flow
blocks may be needed

Hash Table Performance

- Fixed number of buckets
- Excellent, if no overflow blocks
- Degrades considerably when there are many overflow blocks.
 - Might need to go through a chain of overflow blocks

Can improve this by allowing the number of buckets to grow


Outline

- Hash Tables
 - ✓ Secondary storage HT
- Extensible HT
- Linear HT

Extensible Hash Table

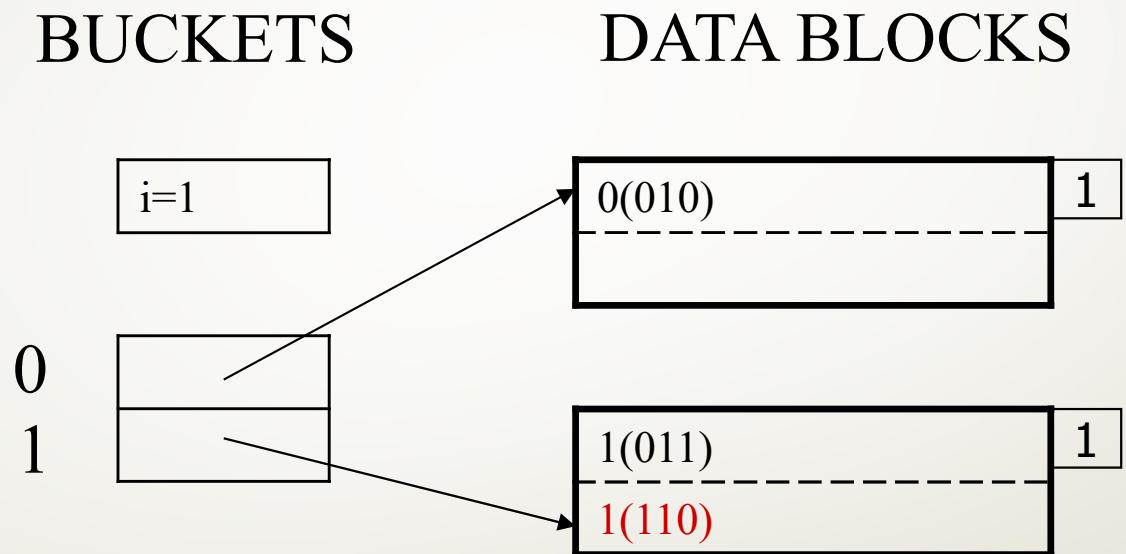
- Array of pointers to blocks instead of array of blocks
- Size of array is allowed to grow. 2x size when it grows
- Don't need a block per bucket. Sparse buckets share a block
- Hash function returns k-bit integers (e.g., k=32)
 - Only use the first $i \ll k$ bits to determine bucket
 - Number of buckets = 2^i

Bit counter on each block indicates how much bits are used for that block



Insertion in Extensible Hash Table

- Insert 1110




Insertion in Extensible Hash Table

- Now insert 1010

BUCKETS


DATA BLOCKS



- Need to split block and extend bucket array
- i becomes 2: done in two steps


Insertion in Extensible Hash Table

Step 1: Extend the buckets



Insertion in Extensible Hash Table


Step 2: Now try to insert 1010



Insertion in Extensible Hash Table

- Now insert 0000: where would it go? Then 0101?
- Need to split block, but not bucket array


BUCKETS DATA BLOCKS



Insertion in Extensible Hash Table


- Now insert 0000: where would it go? Then 0101?
- Need to split block, but not bucket array

BUCKETS DATA BLOCKS



Insertion in Extensible Hash Table

- Now insert 0000: where would it go? Then 0101?
- Need to split block, but not bucket array



Performance: Extensible Hash Table

- No overflow blocks: access always one read for distinct keys
- BUT:
 - Extensions can be **costly and disruptive**
 - After an extension bucket table **may no longer fit in memory**
 - Imagine three records whose keys share the first 20 bits. These three records cannot be in same block (assume two records per block). But a block split would require setting $i = 20$, i.e., accommodating for $2^{20} = 1 \text{ million buckets}$, even though there may be only a few hundred records.

Outline

- Hash Tables
 - ✓ Secondary storage HT
 - ✓ Extensible HT
- Linear HT

Linear Hash Table

- Idea 1: add only one bucket at a time

Problem: n = no longer a power of 2


- Let i be # bits necessary to address n buckets.
 - $i = \text{ceil}(\log_2 n)$
- After computing $h(k)$, use *last* i bits:
 - If last i bits represent a number (say m) $< n$, store the key in bucket m
 - If $m \geq n$, change msb from 1 to 0 (get a number $< n$)
- Idea 2: allow overflow blocks (not expensive to overflow)
- Convention: Read from the right (as opposed to the left)

Linear Hash Table Example

- $N=3 \leq 2^2 = 4$
 - Therefore, only buckets until 10

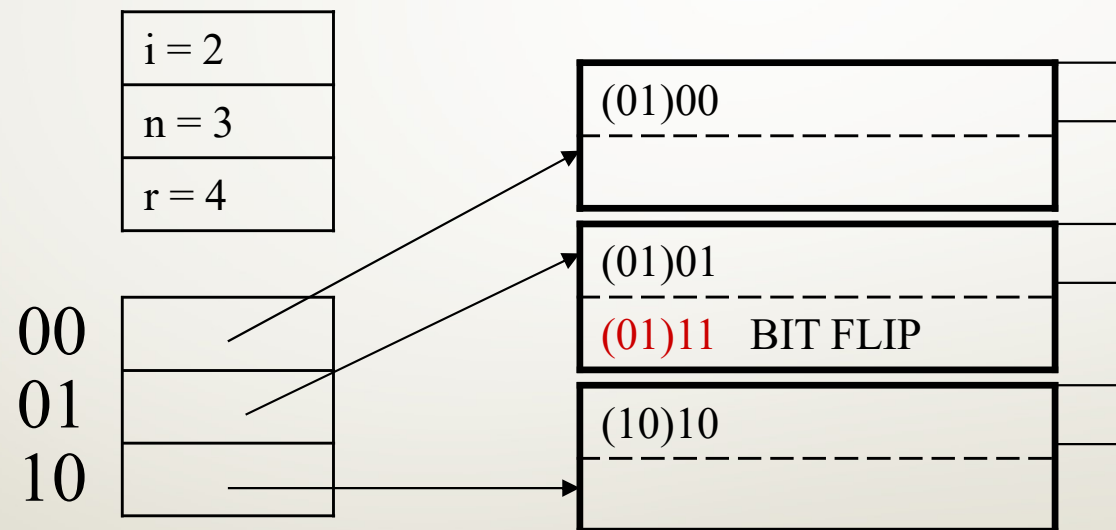
Try to insert 0111

11 is flipped \Rightarrow 01




Linear Hash Table Example

- After inserting 0111



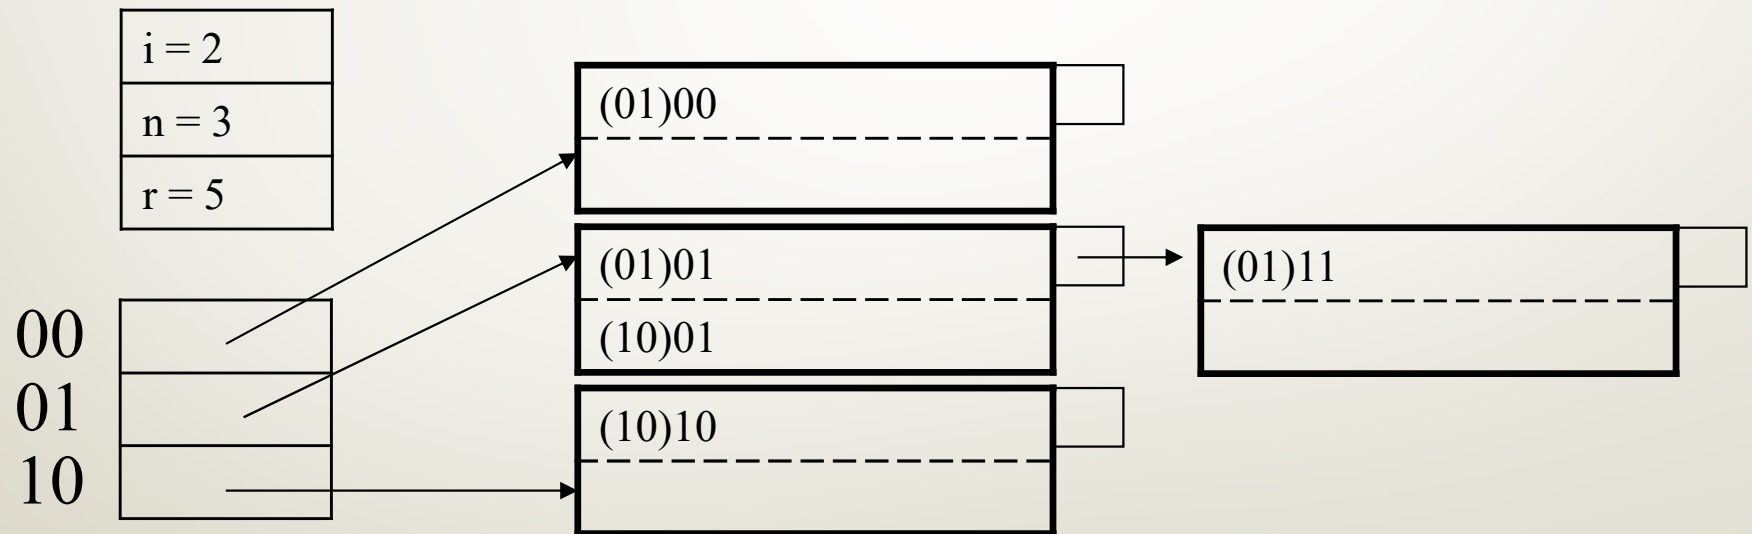
Linear Hash Table Example

- Insert 1001:



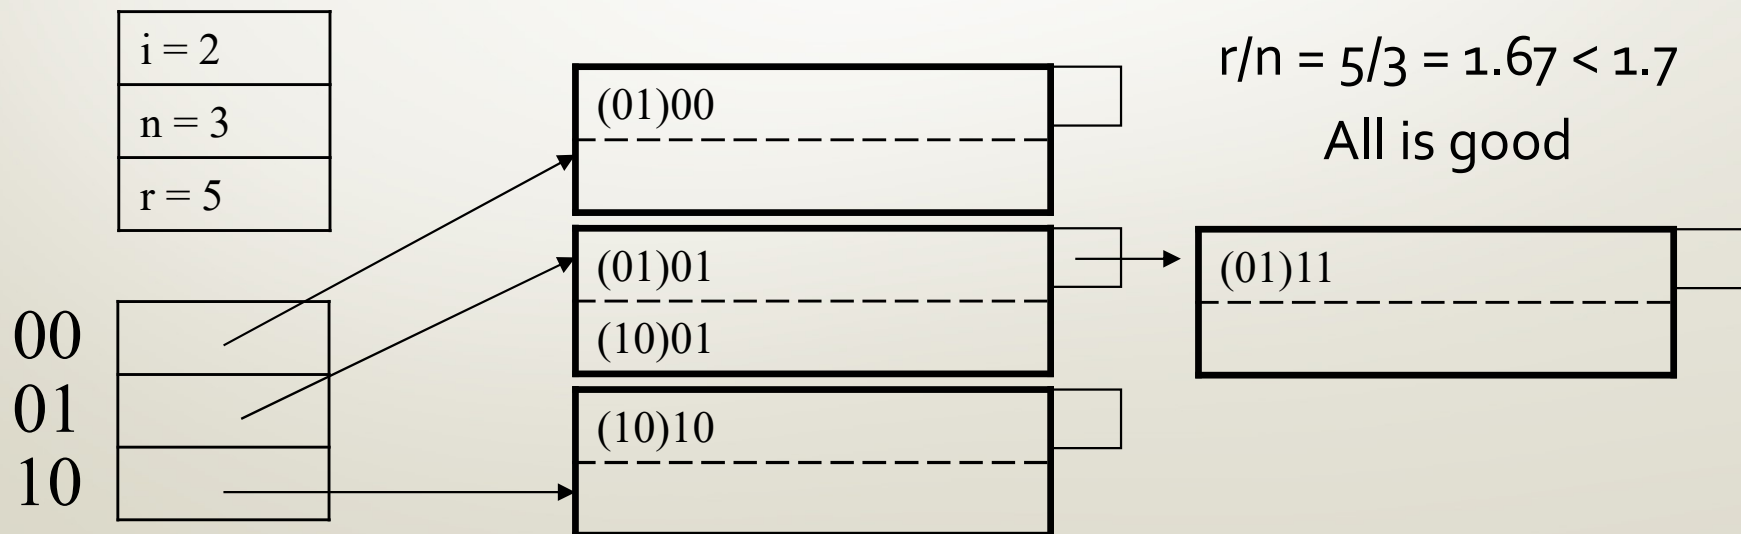
Linear Hash Table Example

- Insert 1001: overflow blocks...



Linear Hash Tables

- Extend $n \rightarrow n+1$ when average number of records per bucket exceeds (say) 85% of total number of records per block
 - e.g., $r/n \leq 0.85 * 2 = 1.7$ (for block size = 2)
- Until then, use overflow blocks (cheaper than adding buckets)

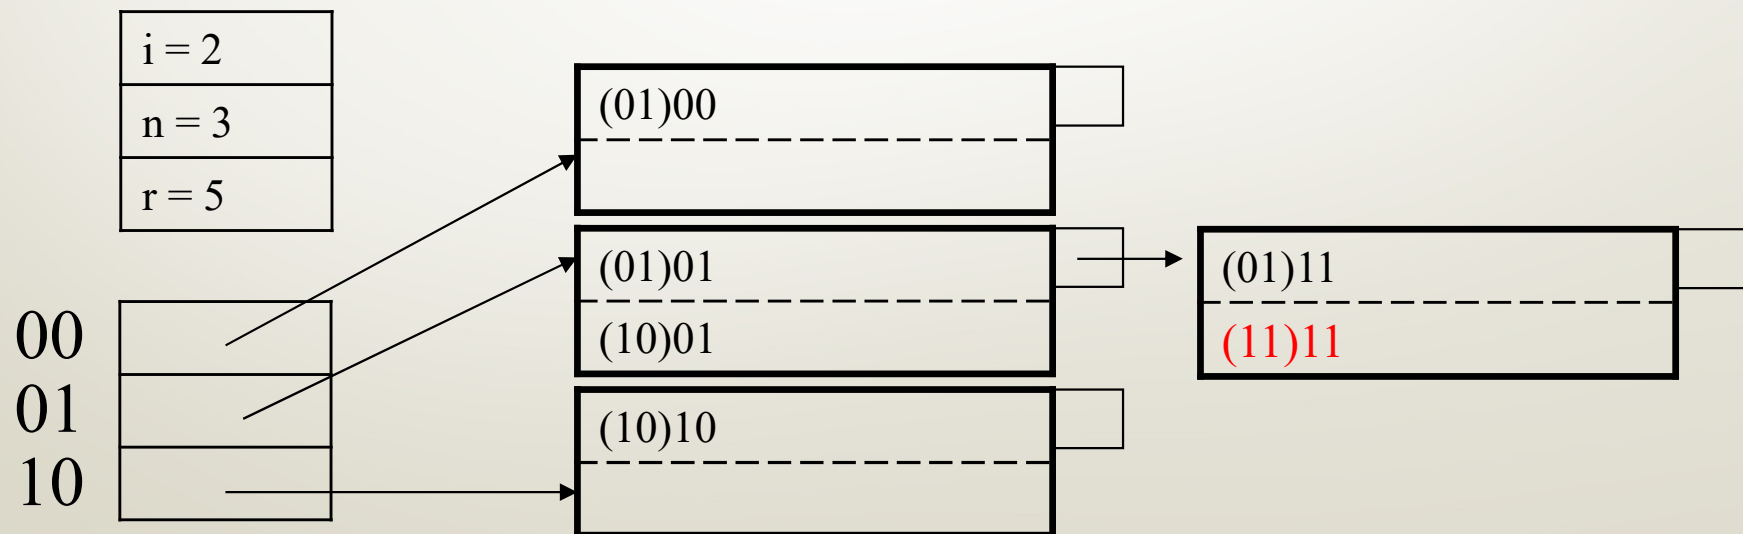


Linear Hash Tables

- Try to insert 1111


$$r/n = 6/3 = 2 > 1.7$$

→ Time to add a bucket




Linear Hash Table Extension

- From $n=3$ to $n=4$




- Only need to touch one block (which one ?)



Linear Hash Table Extension

- From $n=3$ to $n=4$ finished

$$r/n = 6/4 = 1.5 < 1.7 \quad \checkmark$$



Indexing Summary

- B+ Trees (search, insertion, deletion)
 - Good for point and range queries
 - Log time lookup, insertion and deletion because of balanced tree
- Hash Tables (search, insertion)
 - Static hash tables: one I/O lookup, unless long chain of overflow
 - Extensible hash tables: one I/O lookup, extension can take long
 - Linear hash tables: ~ one I/O lookup, cheaper extension
- No panacea; dependent on data and use case